

Question 1 :

```
DROP TABLE if exists TIME CASCADE;
```

```
CREATE TABLE time (  
    TimeId serial primary key,  
    OrderDate date NOT NULL,  
    DayOfWeek character(10) NOT NULL,  
    Month character(10) NOT NULL,  
    Year integer NOT NULL  
);
```

```
INSERT INTO time ( OrderDate, DayOfWeek, Month ,Year )  
SELECT DISTINCT cust_order.orderdate ,  
    to_char( cust_order.orderdate , 'Day') as DayOfWeeek , to_char( cust_order.orderdate  
, 'Month' ) as Month,  
EXTRACT(Year FROM cust_order.orderdate ) as Year  
FROM cust_order order by orderdate asc ;
```

```
SELECT COUNT(*) FROM TIME;
```

```
drop materialized view if exists Sales CASCADE ;  
CREATE materialized view Sales as  
select customer.CustomerId as CustomerId , TIME.TimeId as TimeId , book.isbn AS ISBN ,  
SUM(order_detail.quantity * book.price) AS Amnt  
from book NATURAL join order_detail NATURAL JOIN cust_order NATURAL JOIN customer  
NATURAL JOIN TIME  
group by customer.CustomerId , TIME.TimeId , book.isbn  
order by customer.CustomerId , TIME.TimeId , book.isbn ;
```

```
select COUNT(*) from Sales;
```

Question 2

```
create materialized view avg_amnt_view as select customerid,  
avg(amnt) as avg_amnt from sales group by customerid;  
select avg(avg_amnt) from avg_amnt_view;
```

```
select avg(amnt) from sales;
```

```
wusong3=> create materialized view avg_amnt_view as select customerid,
```

```
wusong3-> avg(amnt) as avg_amnt from sales group by customerid;
```

```
SELECT 104
```

```
wusong3=> select avg(avg_amnt) from avg_amnt_view;
```

```
avg
```

```
-----
```

```
202.9588687852809865
```

```
(1 row)
```

```
wusong3=> select avg(amnt) from sales;
```

```
avg
```

```
-----
```

```
161.3691588785046729
```

```
(1 row)
```

The second one is the correct answer because average of average is ignoring the weight of each column.

Question 3 .

a)

```
SELECT SUM (AMNT) AS money, customer.customerid , customer.l_name , customer.f_name
```

```
from sales NATURAL join customer
```

```
group by customer.customerid
```

```
ORDER BY money DESC limit 5;
```

```
wusong3=> SELECT SUM (AMNT) AS money, customer.customerid , customer.l_name ,
```

```
customer.f_name from sales NATURAL join customer
```

```
wusong3-> group by customer.customerid , customer.l_name , customer.f_name
```

```
wusong3-> ORDER BY money DESC limit 5;
```

```
money | customerid | l_name | f_name
```

```
-----+-----+-----+-----
```

17810.00	1	Jacson	Kirk
14100.00	3	Andree	Peter
11780.00	14	Anslow	Craig
7145.00	2	Leow	May-N
6095.00	79	Liang	Jiajun

```
(5 rows)
```

roll up and rank are used.

b)

rank is used .

```
drop materialized view if exists BestCS CASCADE;
create materialized view BestCS AS
  SELECT SUM (AMNT) AS money, customer.customerid AS CID, customer.l_name ,
customer.f_name from sales NATURAL join customer
  group by customer.customerid ,customer.l_name , customer.f_name
  ORDER BY money DESC limit 1 ;
```

```
drop materialized view if exists BestCSorders CASCADE;
create materialized view BestCSorders AS
  SELECT order_detail.orderid, sum(order_detail.quantity * book.price) AS BST
  FROM order_detail NATURAL JOIN book NATURAL JOIN cust_order NATURAL JOIN
customer
  WHERE cust_order.customerid = (SELECT CID FROM BestCS )
  GROUP BY orderid;
```

```
drop materialized view if exists ord_avg_amntV CASCADE;
create materialized view ord_avg_amntV AS
  select avg(costForEachOrder.oa) as ord_avg_amnt from
  (SELECT order_detail.orderid, sum(order_detail.quantity * book.price) AS oa
  FROM order_detail NATURAL JOIN book
  GROUP BY orderid) as costForEachOrder;
```

```
drop materialized view if exists numerator CASCADE ;
create materialized view numerator AS
SELECT COUNT(*) FROM BestCSorders , ord_avg_amntV WHERE BST >
ord_avg_amntV.ord_avg_amnt ;
```

```
drop materialized view if exists denominator CASCADE;
create materialized view denominator AS
```

```
SELECT COUNT(cust_order.orderid) as no_of_ord from customer NATURAL join cust_order
where customer.customerid = (SELECT CID FROM BestCS ) ;
```

```
drop materialized view if exists division CASCADE;
create materialized view division AS
SELECT COUNT as numerator, no_of_ord as denominator from numerator NATURAL join
denominator ;
```

```
select numerator/denominator::float as perc_of_ord from division ;
```

```
wusong3=> select numerator/denominator::float as perc_of_ord from division ;
          perc_of_ord
```

```
-----
0.714285714285714
(1 row)
```

since perc_of_ord is 71% , we estimate that "e that the best buyer has issued a greater (than average) to medium number of orders with greater (than average) amounts of money,"

Question 4 .

a)

```
drop materialized view if exists View1 CASCADE ;
CREATE MATERIALIZED VIEW View1 AS
SELECT c.CustomerId, F_Name, L_Name, District, TimeId,
DayOfWeek, ISBN, Amnt
FROM Sales NATURAL JOIN Customer c NATURAL JOIN Time;
```

```
drop materialized view if exists View2 CASCADE;
CREATE MATERIALIZED VIEW View2 AS
SELECT c.CustomerId, F_Name, L_Name, Year, SUM(Amnt)
FROM Sales NATURAL JOIN Customer c NATURAL JOIN Time
GROUP BY c.CustomerId, F_Name, L_Name, Year;
```

--1. The "Book Orders Database"
EXPLAIN ANALYZE

```

SELECT SUM (AMNT) AS money, customer.customerid , customer.l_name , customer.f_name
from
(select customer.CustomerId as CustomerId ,TIME.Timeid as Timeid , book.isbn AS ISBN ,
SUM(order_detail.quantity * book.price) AS Amnt
from book NATURAL join order_detail NATURAL JOIN cust_order NATURAL JOIN customer
NATURAL JOIN TIME
group by customer.CustomerId , TIME.Timeid ,book.isbn
order by customer.CustomerId , TIME.Timeid , book.isbn ) as sales
NATURAL join customer
group by customer.customerid , customer.l_name , customer.f_name
ORDER BY money DESC limit 5;
VACUUM (ANALYZE) ;

```

```

wusong3=> EXPLAIN ANALYZE
wusong3-> SELECT SUM (AMNT) AS money, customer.customerid , customer.l_name ,
customer.f_name from
wusong3-> (select customer.CustomerId as CustomerId ,TIME.Timeid as Timeid , book.isbn
AS ISBN , SUM(order_detail.quantity * book.price) AS Amnt
wusong3(> from book NATURAL join order_detail NATURAL JOIN cust_order NATURAL JOIN
customer NATURAL JOIN TIME
wusong3(> group by customer.CustomerId , TIME.Timeid ,book.isbn
wusong3(> order by customer.CustomerId , TIME.Timeid , book.isbn ) as sales
wusong3-> NATURAL join customer
wusong3-> group by customer.customerid , customer.l_name , customer.f_name
wusong3-> ORDER BY money DESC limit 5;

```

QUERY PLAN

```

-----
Limit (cost=197.83..197.84 rows=5 width=78) (actual time=24.364..24.380 rows=5 loops=1)
-> Sort (cost=197.83..198.12 rows=118 width=78) (actual time=24.359..24.364 rows=5
loops=1)
    Sort Key: (sum((sum(((order_detail.quantity)::numeric * book.price))))
    Sort Method: top-N heapsort  Memory: 25kB
    -> HashAggregate (cost=194.39..195.87 rows=118 width=78) (actual
time=23.991..24.146 rows=104 loops=1)
        Group Key: customer.customerid, customer.l_name, customer.f_name
        -> Hash Join (cost=133.29..187.90 rows=649 width=78) (actual time=12.429..21.238
rows=1070 loops=1)
            Hash Cond: (customer_1.customerid = customer.customerid)
            -> GroupAggregate (cost=128.63..161.63 rows=1100 width=19) (actual
time=12.060..17.896 rows=1070 loops=1)
                Group Key: customer_1.customerid, "time".timeid, book.isbn
                -> Sort (cost=128.63..131.38 rows=1100 width=19) (actual
time=12.033..13.433 rows=1100 loops=1)

```

Sort Key: customer_1.customerid, "time".timeid, book.isbn
Sort Method: quicksort Memory: 134kB
-> Hash Join (cost=23.82..73.06 rows=1100 width=19) (actual time=2.594..9.759 rows=1100 loops=1)
Hash Cond: (order_detail.isbn = book.isbn)
-> Hash Join (cost=22.55..56.67 rows=1100 width=14) (actual time=2.541..6.918 rows=1100 loops=1)
Hash Cond: (order_detail.orderid = cust_order.orderid)
-> Seq Scan on order_detail (cost=0.00..19.00 rows=1100 width=10) (actual time=0.004..1.330 rows=1100 loops=1)
-> Hash (cost=19.77..19.77 rows=222 width=12) (actual time=2.524..2.524 rows=222 loops=1)
Buckets: 1024 Batches: 1 Memory Usage: 10kB
-> Hash Join (cost=9.45..19.77 rows=222 width=12) (actual time=0.709..2.191 rows=222 loops=1)
Hash Cond: (cust_order.orderdate = "time".orderdate)
-> Hash Join (cost=4.65..11.93 rows=222 width=12) (actual time=0.345..1.223 rows=222 loops=1)
Hash Cond: (cust_order.customerid = customer_1.customerid)
-> Seq Scan on cust_order (cost=0.00..4.22 rows=222 width=12) (actual time=0.004..0.275 rows=222 loops=1)
-> Hash (cost=3.18..3.18 rows=118 width=4) (actual time=0.326..0.326 rows=118 loops=1)
Buckets: 1024 Batches: 1 Memory Usage: 5kB
-> Seq Scan on customer customer_1 (cost=0.00..3.18 rows=118 width=4) (actual time=0.003..0.153 rows=118 loops=1)
-> Hash (cost=3.24..3.24 rows=124 width=8) (actual time=0.351..0.351 rows=124 loops=1)
Buckets: 1024 Batches: 1 Memory Usage: 5kB
-> Seq Scan on "time" (cost=0.00..3.24 rows=124 width=8) (actual time=0.006..0.173 rows=124 loops=1)
-> Hash (cost=1.12..1.12 rows=12 width=9) (actual time=0.042..0.042 rows=12 loops=1)
Buckets: 1024 Batches: 1 Memory Usage: 1kB
-> Seq Scan on book (cost=0.00..1.12 rows=12 width=9) (actual time=0.003..0.021 rows=12 loops=1)
-> Hash (cost=3.18..3.18 rows=118 width=46) (actual time=0.354..0.354 rows=118 loops=1)
Buckets: 1024 Batches: 1 Memory Usage: 9kB
-> Seq Scan on customer (cost=0.00..3.18 rows=118 width=46) (actual time=0.007..0.171 rows=118 loops=1)
Planning time: 1.608 ms
Execution time: 24.567 ms

(39 rows)

```
wusong3=> VACUUM (ANALYZE) ;  
WARNING: skipping "pg_authid" --- only superuser can vacuum it  
WARNING: skipping "pg_database" --- only superuser can vacuum it  
WARNING: skipping "pg_db_role_setting" --- only superuser can vacuum it  
WARNING: skipping "pg_tablespace" --- only superuser can vacuum it  
WARNING: skipping "pg_pltemplate" --- only superuser can vacuum it  
WARNING: skipping "pg_auth_members" --- only superuser can vacuum it  
WARNING: skipping "pg_shdepend" --- only superuser can vacuum it  
WARNING: skipping "pg_shdescription" --- only superuser can vacuum it  
WARNING: skipping "pg_shseclabel" --- only superuser can vacuum it
```

VACUUM

```
--2. The Data Mart,  
EXPLAIN ANALYZE  
SELECT SUM (AMNT) AS money, customer.customerid , customer.l_name , customer.f_name  
from sales NATURAL join customer  
group by customer.customerid , customer.l_name , customer.f_name  
ORDER BY money DESC limit 5;  
VACUUM (ANALYZE) ;
```

```
wusong3=> EXPLAIN ANALYZE  
wusong3-> SELECT SUM (AMNT) AS money, customer.customerid , customer.l_name ,  
customer.f_name from sales NATURAL join customer  
wusong3-> group by customer.customerid , customer.l_name , customer.f_name  
wusong3-> ORDER BY money DESC limit 5;
```

QUERY PLAN

```
-----  
-----  
Limit (cost=51.20..51.21 rows=5 width=51) (actual time=7.827..7.842 rows=5 loops=1)  
-> Sort (cost=51.20..51.50 rows=118 width=51) (actual time=7.823..7.828 rows=5 loops=1)  
    Sort Key: (sum(sales.amnt))  
    Sort Method: top-N heapsort  Memory: 25kB  
-> HashAggregate (cost=47.77..49.24 rows=118 width=51) (actual time=7.451..7.611  
rows=104 loops=1)  
    Group Key: customer.customerid, customer.l_name, customer.f_name  
-> Hash Join (cost=4.65..37.07 rows=1070 width=51) (actual time=0.381..4.743  
rows=1070 loops=1)  
    Hash Cond: (sales.customerid = customer.customerid)  
-> Seq Scan on sales (cost=0.00..17.70 rows=1070 width=9) (actual  
time=0.008..1.361 rows=1070 loops=1)
```

-> Hash (cost=3.18..3.18 rows=118 width=46) (actual time=0.355..0.355 rows=118 loops=1)
Buckets: 1024 Batches: 1 Memory Usage: 9kB
-> Seq Scan on customer (cost=0.00..3.18 rows=118 width=46) (actual time=0.004..0.162 rows=118 loops=1)
Planning time: 0.420 ms
Execution time: 7.948 ms
(14 rows)

wusong3=> VACUUM (ANALYZE) ;
WARNING: skipping "pg_authid" --- only superuser can vacuum it
WARNING: skipping "pg_database" --- only superuser can vacuum it
WARNING: skipping "pg_db_role_setting" --- only superuser can vacuum it
WARNING: skipping "pg_tablespace" --- only superuser can vacuum it
WARNING: skipping "pg_pltemplate" --- only superuser can vacuum it
WARNING: skipping "pg_auth_members" --- only superuser can vacuum it
WARNING: skipping "pg_shdepend" --- only superuser can vacuum it
WARNING: skipping "pg_shdescription" --- only superuser can vacuum it
WARNING: skipping "pg_shseclabel" --- only superuser can vacuum it
VACUUM

--3. The view View1,

EXPLAIN ANALYZE
SELECT SUM (AMNT) AS money FROM View1 GROUP BY customerId
ORDER BY money DESC limit 5;
VACUUM ANALYZE ;

wusong3=> EXPLAIN ANALYZE
wusong3-> SELECT SUM (AMNT) AS money FROM View1 GROUP BY customerId
wusong3-> ORDER BY money DESC limit 5;

QUERY PLAN

Limit (cost=36.08..36.09 rows=5 width=9) (actual time=3.829..3.845 rows=5 loops=1)
-> Sort (cost=36.08..36.34 rows=104 width=9) (actual time=3.824..3.828 rows=5 loops=1)
Sort Key: (sum(amnt))
Sort Method: top-N heapsort Memory: 25kB
-> HashAggregate (cost=33.05..34.35 rows=104 width=9) (actual time=3.452..3.622 rows=104 loops=1)
Group Key: customerId
-> Seq Scan on view1 (cost=0.00..27.70 rows=1070 width=9) (actual time=0.006..1.344 rows=1070 loops=1)

Planning time: 0.213 ms
Execution time: 3.924 ms
(9 rows)

```
wusong3=> VACUUM ANALYZE ;  
WARNING: skipping "pg_authid" --- only superuser can vacuum it  
WARNING: skipping "pg_database" --- only superuser can vacuum it  
WARNING: skipping "pg_db_role_setting" --- only superuser can vacuum it  
WARNING: skipping "pg_tablespace" --- only superuser can vacuum it  
WARNING: skipping "pg_pltemplate" --- only superuser can vacuum it  
WARNING: skipping "pg_auth_members" --- only superuser can vacuum it  
WARNING: skipping "pg_shdepend" --- only superuser can vacuum it  
WARNING: skipping "pg_shdescription" --- only superuser can vacuum it  
WARNING: skipping "pg_shseclabel" --- only superuser can vacuum it  
VACUUM
```

--4. The view View2.

```
EXPLAIN ANALYZE  
SELECT sum(sum) AS money , customerid , f_name , l_name FROM View2  
group by customerid , f_name , l_name  
ORDER BY money DESC limit 5 ;  
VACUUM ANALYZE ;
```

```
wusong3=> EXPLAIN ANALYZE  
wusong3-> SELECT sum(sum) AS money , customerid , f_name , l_name FROM View2  
wusong3-> group by customerid , f_name , l_name  
wusong3-> ORDER BY money DESC limit 5 ;  
QUERY PLAN
```

```
-----  
Limit (cost=7.67..7.68 rows=5 width=51) (actual time=0.928..0.944 rows=5 loops=1)  
-> Sort (cost=7.67..7.93 rows=104 width=51) (actual time=0.925..0.930 rows=5 loops=1)  
    Sort Key: (sum(sum))  
    Sort Method: top-N heapsort Memory: 25kB  
-> HashAggregate (cost=4.64..5.94 rows=104 width=51) (actual time=0.551..0.703  
rows=104 loops=1)  
    Group Key: customerid, f_name, l_name  
-> Seq Scan on view2 (cost=0.00..3.32 rows=132 width=51) (actual  
time=0.011..0.193 rows=132 loops=1)  
Planning time: 0.144 ms  
Execution time: 1.025 ms  
(9 rows)
```

```
wusong3=> VACUUM ANALYZE ;
WARNING: skipping "pg_authid" --- only superuser can vacuum it
WARNING: skipping "pg_database" --- only superuser can vacuum it
WARNING: skipping "pg_db_role_setting" --- only superuser can vacuum it
WARNING: skipping "pg_tablespace" --- only superuser can vacuum it
WARNING: skipping "pg_pltemplate" --- only superuser can vacuum it
WARNING: skipping "pg_auth_members" --- only superuser can vacuum it
WARNING: skipping "pg_shdepend" --- only superuser can vacuum it
WARNING: skipping "pg_shdescription" --- only superuser can vacuum it
WARNING: skipping "pg_shseclabel" --- only superuser can vacuum it
VACUUM
```

Explain the findings: the cost of the execution of the QUERY drops as more aggregate function perform in MATERIALIZED view level .

b)

```
drop materialized view if exists View3 CASCADE;
CREATE MATERIALIZED VIEW View3 AS
SELECT District, TimeId, DayOfWeek, ISBN, SUM(Amnt)
FROM Sales NATURAL JOIN Customer NATURAL JOIN Time
GROUP BY District, TimeId, DayOfWeek, ISBN;
```

--1. The "Book Orders Database"

```
EXPLAIN ANALYZE
SELECT SUM(Amnt) , country FROM
customer NATURAL JOIN
(select customer.CustomerId as CustomerId ,TIME.TimeId as TimeId , book.isbn AS ISBN ,
SUM(order_detail.quantity * book.price) AS Amnt
from book NATURAL join order_detail NATURAL JOIN cust_order NATURAL JOIN customer
NATURAL JOIN TIME
group by customer.CustomerId , TIME.TimeId ,book.isbn
order by customer.CustomerId , TIME.TimeId , book.isbn ) AS sales
GROUP BY country order BY SUM DESC limit 1;
VACUUM ANALYZE ;
```

```
wusong3=> EXPLAIN ANALYZE
wusong3-> SELECT SUM(Amnt) , country FROM
wusong3-> customer NATURAL JOIN
```

```
wusong3-> (select customer.CustomerId as CustomerId ,TIME.TimeId as TimeId , book.isbn
AS ISBN , SUM(order_detail.quantity * book.price) AS Amnt
wusong3(> from book NATURAL join order_detail NATURAL JOIN cust_order NATURAL JOIN
customer NATURAL JOIN TIME
wusong3(> group by customer.CustomerId , TIME.TimeId ,book.isbn
wusong3(> order by customer.CustomerId , TIME.TimeId , book.isbn ) AS sales
wusong3-> GROUP BY country order BY SUM DESC limit 1;
```

QUERY PLAN

```
-----
Limit (cost=191.27..191.27 rows=1 width=48) (actual time=23.613..23.615 rows=1 loops=1)
  -> Sort (cost=191.27..191.29 rows=7 width=48) (actual time=23.608..23.608 rows=1
loops=1)
    Sort Key: (sum((sum(((order_detail.quantity)::numeric * book.price))))
    Sort Method: top-N heapsort Memory: 25kB
    -> HashAggregate (cost=191.15..191.24 rows=7 width=48) (actual time=23.540..23.552
rows=7 loops=1)
      Group Key: customer.country
      -> Hash Join (cost=133.29..187.90 rows=649 width=48) (actual time=12.594..21.422
rows=1070 loops=1)
        Hash Cond: (customer_1.customerid = customer.customerid)
        -> GroupAggregate (cost=128.63..161.63 rows=1100 width=28) (actual
time=12.226..18.079 rows=1070 loops=1)
          Group Key: customer_1.customerid, "time".timeid, book.isbn
          -> Sort (cost=128.63..131.38 rows=1100 width=28) (actual
time=12.200..13.565 rows=1100 loops=1)
            Sort Key: customer_1.customerid, "time".timeid, book.isbn
            Sort Method: quicksort Memory: 134kB
            -> Hash Join (cost=23.82..73.06 rows=1100 width=28) (actual
time=2.617..9.985 rows=1100 loops=1)
              Hash Cond: (order_detail.isbn = book.isbn)
              -> Hash Join (cost=22.55..56.67 rows=1100 width=14) (actual
time=2.563..7.010 rows=1100 loops=1)
                Hash Cond: (order_detail.orderid = cust_order.orderid)
                -> Seq Scan on order_detail (cost=0.00..19.00 rows=1100
width=10) (actual time=0.008..1.410 rows=1100 loops=1)
                -> Hash (cost=19.77..19.77 rows=222 width=12) (actual
time=2.538..2.538 rows=222 loops=1)
                  Buckets: 1024 Batches: 1 Memory Usage: 10kB
                  -> Hash Join (cost=9.45..19.77 rows=222 width=12) (actual
time=0.712..2.198 rows=222 loops=1)
                    Hash Cond: (cust_order.orderdate = "time".orderdate)
                    -> Hash Join (cost=4.65..11.93 rows=222 width=12)
(actual time=0.350..1.221 rows=222 loops=1)
```

```

Hash Cond: (cust_order.customerid =
customer_1.customerid)
-> Seq Scan on cust_order (cost=0.00..4.22 rows=222
width=12) (actual time=0.008..0.279 rows=222 loops=1)
-> Hash (cost=3.18..3.18 rows=118 width=4) (actual
time=0.327..0.327 rows=118 loops=1)
Buckets: 1024 Batches: 1 Memory Usage: 5kB
-> Seq Scan on customer customer_1
(cost=0.00..3.18 rows=118 width=4) (actual time=0.005..0.157 rows=118 loops=1)
-> Hash (cost=3.24..3.24 rows=124 width=8) (actual
time=0.350..0.350 rows=124 loops=1)
Buckets: 1024 Batches: 1 Memory Usage: 5kB
-> Seq Scan on "time" (cost=0.00..3.24 rows=124
width=8) (actual time=0.009..0.170 rows=124 loops=1)
-> Hash (cost=1.12..1.12 rows=12 width=18) (actual
time=0.042..0.042 rows=12 loops=1)
Buckets: 1024 Batches: 1 Memory Usage: 1kB
-> Seq Scan on book (cost=0.00..1.12 rows=12 width=18) (actual
time=0.005..0.019 rows=12 loops=1)
-> Hash (cost=3.18..3.18 rows=118 width=20) (actual time=0.352..0.352
rows=118 loops=1)
Buckets: 1024 Batches: 1 Memory Usage: 6kB
-> Seq Scan on customer (cost=0.00..3.18 rows=118 width=20) (actual
time=0.009..0.174 rows=118 loops=1)
Planning time: 1.400 ms
Execution time: 23.793 ms
(39 rows)

```

```

wusong3=> VACUUM ANALYZE ;
WARNING: skipping "pg_authid" --- only superuser can vacuum it
WARNING: skipping "pg_database" --- only superuser can vacuum it
WARNING: skipping "pg_db_role_setting" --- only superuser can vacuum it
WARNING: skipping "pg_tablespace" --- only superuser can vacuum it
WARNING: skipping "pg_pltemplate" --- only superuser can vacuum it
WARNING: skipping "pg_auth_members" --- only superuser can vacuum it
WARNING: skipping "pg_shdepend" --- only superuser can vacuum it
WARNING: skipping "pg_shdescription" --- only superuser can vacuum it
WARNING: skipping "pg_shseclabel" --- only superuser can vacuum it
VACUUM

```

```

--2. The Data Mart,
EXPLAIN ANALYZE
SELECT SUM(Amnt) , country FROM

```

```
customer NATURAL JOIN
sales
GROUP BY country order BY SUM DESC limit 1;
VACUUM ANALYZE ;
```

```
wusong3=> EXPLAIN ANALYZE
wusong3-> SELECT SUM(Amnt) , country FROM
wusong3-> customer NATURAL JOIN
wusong3-> sales
wusong3-> GROUP BY country order BY SUM DESC limit 1;
QUERY PLAN
```

```
-----
Limit (cost=42.54..42.54 rows=1 width=21) (actual time=6.780..6.782 rows=1 loops=1)
  -> Sort (cost=42.54..42.56 rows=7 width=21) (actual time=6.776..6.776 rows=1 loops=1)
        Sort Key: (sum(sales.amnt))
        Sort Method: top-N heapsort Memory: 25kB
        -> HashAggregate (cost=42.42..42.50 rows=7 width=21) (actual time=6.736..6.745
rows=7 loops=1)
            Group Key: customer.country
            -> Hash Join (cost=4.65..37.07 rows=1070 width=21) (actual time=0.376..4.667
rows=1070 loops=1)
                Hash Cond: (sales.customerid = customer.customerid)
                -> Seq Scan on sales (cost=0.00..17.70 rows=1070 width=9) (actual
time=0.008..1.327 rows=1070 loops=1)
                -> Hash (cost=3.18..3.18 rows=118 width=20) (actual time=0.350..0.350
rows=118 loops=1)
                    Buckets: 1024 Batches: 1 Memory Usage: 6kB
                    -> Seq Scan on customer (cost=0.00..3.18 rows=118 width=20) (actual
time=0.004..0.169 rows=118 loops=1)
Planning time: 0.473 ms
Execution time: 6.871 ms
(14 rows)
```

```
wusong3=> VACUUM ANALYZE ;
WARNING: skipping "pg_authid" --- only superuser can vacuum it
WARNING: skipping "pg_database" --- only superuser can vacuum it
WARNING: skipping "pg_db_role_setting" --- only superuser can vacuum it
WARNING: skipping "pg_tablespace" --- only superuser can vacuum it
WARNING: skipping "pg_pltemplate" --- only superuser can vacuum it
WARNING: skipping "pg_auth_members" --- only superuser can vacuum it
WARNING: skipping "pg_shdepend" --- only superuser can vacuum it
WARNING: skipping "pg_shdescription" --- only superuser can vacuum it
WARNING: skipping "pg_shseclabel" --- only superuser can vacuum it
```

VACUUM

```
--3. The view View2,  
EXPLAIN ANALYZE  
SELECT SUM(sum) , country FROM  
view2 NATURAL JOIN  
customer  
GROUP BY country order BY SUM DESC limit 1;  
VACUUM ANALYZE ;
```

```
wusong3=> EXPLAIN ANALYZE  
wusong3-> SELECT SUM(sum) , country FROM  
wusong3-> view2 NATURAL JOIN  
wusong3-> customer  
wusong3-> GROUP BY country order BY SUM DESC limit 1;  
QUERY PLAN
```

```
-----  
-----  
-----  
Limit (cost=10.09..10.09 rows=1 width=21) (actual time=1.450..1.451 rows=1 loops=1)  
-> Sort (cost=10.09..10.09 rows=1 width=21) (actual time=1.445..1.445 rows=1 loops=1)  
    Sort Key: (sum(view2.sum))  
    Sort Method: top-N heapsort  Memory: 25kB  
-> HashAggregate (cost=10.06..10.08 rows=1 width=21) (actual time=1.406..1.418  
rows=7 loops=1)  
    Group Key: customer.country  
-> Hash Join (cost=5.25..10.06 rows=1 width=21) (actual time=0.432..1.106  
rows=132 loops=1)  
    Hash Cond: ((view2.customerid = customer.customerid) AND (view2.f_name =  
customer.f_name) AND (view2.l_name = customer.l  
_name))  
-> Seq Scan on view2 (cost=0.00..3.32 rows=132 width=51) (actual  
time=0.007..0.174 rows=132 loops=1)  
-> Hash (cost=3.18..3.18 rows=118 width=62) (actual time=0.400..0.400  
rows=118 loops=1)  
    Buckets: 1024  Batches: 1  Memory Usage: 11kB  
-> Seq Scan on customer (cost=0.00..3.18 rows=118 width=62) (actual  
time=0.005..0.167 rows=118 loops=1)  
Planning time: 0.929 ms  
Execution time: 1.547 ms  
(14 rows)
```

```
wusong3=> VACUUM ANALYZE ;
WARNING: skipping "pg_authid" --- only superuser can vacuum it
WARNING: skipping "pg_database" --- only superuser can vacuum it
WARNING: skipping "pg_db_role_setting" --- only superuser can vacuum it
WARNING: skipping "pg_tablespace" --- only superuser can vacuum it
WARNING: skipping "pg_pltemplate" --- only superuser can vacuum it
WARNING: skipping "pg_auth_members" --- only superuser can vacuum it
WARNING: skipping "pg_shdepend" --- only superuser can vacuum it
WARNING: skipping "pg_shdescription" --- only superuser can vacuum it
WARNING: skipping "pg_shseclabel" --- only superuser can vacuum it
VACUUM
```

--4. The view View3.

```
EXPLAIN ANALYZE
SELECT SUM(sum) , country FROM
view3 NATURAL JOIN
(select distinct district , country from customer) as cc
GROUP BY country order BY SUM DESC limit 1;
VACUUM ANALYZE ;
```

```
wusong3=> EXPLAIN ANALYZE
wusong3-> SELECT SUM(sum) , country FROM
wusong3-> view3 NATURAL JOIN
wusong3-> (select distinct district , country from customer) as cc
wusong3-> GROUP BY country order BY SUM DESC limit 1;
QUERY PLAN
```

```
-----
-----
----
Limit (cost=43.54..43.55 rows=1 width=21) (actual time=6.830..6.832 rows=1 loops=1)
-> Sort (cost=43.54..43.59 rows=17 width=21) (actual time=6.826..6.826 rows=1 loops=1)
    Sort Key: (sum(view3.sum))
    Sort Method: top-N heapsort Memory: 25kB
    -> HashAggregate (cost=43.25..43.46 rows=17 width=21) (actual time=6.786..6.798
rows=7 loops=1)
        Group Key: cc.country
        -> Hash Join (cost=4.32..38.22 rows=1006 width=21) (actual time=0.535..4.849
rows=1006 loops=1)
            Hash Cond: (view3.district = cc.district)
            -> Seq Scan on view3 (cost=0.00..20.06 rows=1006 width=21) (actual
time=0.011..1.261 rows=1006 loops=1)
```

-> Hash (cost=4.11..4.11 rows=17 width=32) (actual time=0.505..0.505 rows=17 loops=1)

Buckets: 1024 Batches: 1 Memory Usage: 2kB

-> Subquery Scan on cc (cost=3.77..4.11 rows=17 width=32) (actual time=0.413..0.475 rows=17 loops=1)

-> HashAggregate (cost=3.77..3.94 rows=17 width=32) (actual time=0.409..0.432 rows=17 loops=1)

Group Key: customer.district, customer.country

-> Seq Scan on customer (cost=0.00..3.18 rows=118 width=32) (actual time=0.007..0.177 rows=118 loops=1)

Planning time: 0.294 ms

Execution time: 6.939 ms

(17 rows)

wusong3=> VACUUM ANALYZE ;

WARNING: skipping "pg_authid" --- only superuser can vacuum it

WARNING: skipping "pg_database" --- only superuser can vacuum it

WARNING: skipping "pg_db_role_setting" --- only superuser can vacuum it

WARNING: skipping "pg_tablespace" --- only superuser can vacuum it

WARNING: skipping "pg_pltemplate" --- only superuser can vacuum it

WARNING: skipping "pg_auth_members" --- only superuser can vacuum it

WARNING: skipping "pg_shdepend" --- only superuser can vacuum it

WARNING: skipping "pg_shdescription" --- only superuser can vacuum it

WARNING: skipping "pg_shseclabel" --- only superuser can vacuum it

VACUUM

Explain the findings: the cost of the execution of the QUERY drops as more aggregate function perform in MATERIALIZED view level .

Question 5 .

a)

```
SELECT DISTINCT * FROM (
  SELECT customerid , f_name , city , sum(amnt) OVER ws, avg(amnt) OVER wa
  FROM sales NATURAL join time NATURAL JOIN customer
  WHERE Month in ('April', 'May' )and Year = 2017
  WINDOW ws AS (PARTITION BY customerid ) ,
         wa AS (PARTITION BY city )
) AS T ORDER BY CITY ;
```


B)

DROP MATERIALIZED VIEW IF EXISTS V1 CASCADE;

CREATE MATERIALIZED VIEW V1 AS

SELECT city, timeid, OrderDate AS day, sum(amnt) AS ST

FROM sales NATURAL JOIN customer NATURAL JOIN time

WHERE Year = 2017 AND Month IN ('April', 'May')

GROUP BY city, timeid, OrderDate

ORDER BY city, timeid , OrderDate;

SELECT city, timeid, day, ST AS "sum(amnt)", sum(ST) OVER W1 AS cumulative_sum

FROM V1 WINDOW W1 AS (PARTITION BY city ORDER BY timeid)

ORDER BY city, timeid , day ,st;