

VICTORIA UNIVERSITY OF WELLINGTON
Te Whare Wananga o te Upoko o te Ika a Maui



Basic OLAP Queries

Lecturer : Dr. Pavle Mogin

SWEN 432
*Advanced Database Design and
Implementation*

Plan for OLAP Queries

- Basic multidimensional OLAP queries:
 - Roll – up and Drill – down
 - Slice and Dice
 - Pivoting

- *Readings :*
 - *Ramakrishnan, Gehrke: “Database Management Systems”, Chapter 25, Sections 25.3, 25.4, and 25.6*
 - *Chaudhuri and Dayal : “An Overview of Datawarehousing and OLAP Technologies”*

Roll – Up and Drill - Down

- ***Roll – up*** operation corresponds to taking the current aggregation level of fact values and doing a further aggregation (getting coarser grained data)
- Each roll-up operation can be expressed using a SQL/92 SELECT...AGG()... GROUP BY... statement
- ***Drill – down*** is just an opposite operation of roll – up, but it requires the access to data of finer granularity

Classification of Roll – Up

- **Roll – up** operations can be classified onto:
 - **Dimensional** roll-ups that are done by dropping one or more dimensions
 - **Hierarchical** roll-ups that are done by climbing up the attribute hierarchies of dimensions
 - **Mixed** roll-ups that combine the previous two techniques
- Note that in an ultimate case, the hierarchical roll-up above the top level of an attribute hierarchy (to the attribute “all”) can be viewed as converting to a dimensional roll-up

Example Star Schema

- In the examples that follow, we shall consider the following star schema:
 - Fact Table :
 - *Sales* (*ShopId*, *ProdId*, *TimeId*, *Total*)
 - *Total* is a fact
 - Dimension tables:
 - *Location*(*ShopId*, *City*),
 - *Product*(*ProdId*, *ProdName*, *Type*) ,
 - *Time*(*TimeId*, *Week*)
 - Dimension attribute hierarchies:
 - *ShopId*→*City* (location hierarchy)
 - *ProdId*→*Type* (product hierarchy)
 - *TimeId*→*Week* (time hierarchy)

Dimensional Roll – Ups

- One dimension dropped:

```
SELECT ProdId, ShopId, SUM(Total) AS Total_by_Prod_Loc
FROM Sales
GROUP BY ProdId, ShopId ;
```

- Two dimensions dropped:

```
SELECT ProdId, SUM(Total) AS Total_by_Prod
FROM Sales
GROUP BY ProdId ;
```

- All dimensions dropped:

```
SELECT SUM(Total) AS Overall_Total
FROM Sales;
```

Hierarchical Roll- Ups (one dimension)

- One roll – up would give sales by product, **city** and day

```
SELECT ProdId, City, TimeId, SUM(Total) AS City_Total  
FROM Sales NATURAL JOIN Location  
GROUP BY City, ProdId, TimeId;
```
- The second roll – up would give sales by **product type**, shop and day

```
SELECT Type, ShopId, TimeId, SUM(Total) AS Type_Total  
FROM Sales NATURAL JOIN Product  
GROUP BY Type, ShopId, TimeId;
```
- The third roll – up **Week_Total** would give sales by product, shop and **week**

```
SELECT ProductId, ShopId, Week, SUM(Total) AS Week_Total  
FROM Sales NATURAL JOIN Time  
GROUP BY Week, ShopId, ProductId;
```

Hierarchical Roll-Ups (2d and 3d)

- Two dimensional roll-up:

```
SELECT Type, City, TimeId, SUM(Total) AS Type_City_Tot  
FROM Sales NATURAL JOIN Product NATURAL JOIN Location  
GROUP BY Type, City, TimeId ;
```

(There would be also two other 2d roll-ups possible (***Type_Week*** and ***City_Week***) but these are not shown)

- Three dimensional roll-up:

```
SELECT Type, City, Week, SUM(Total) AS Type_City_Week_Tot  
FROM Sales NATURAL JOIN Product NATURAL JOIN Location  
NATURAL JOIN Time  
GROUP BY Type, City, Week ;
```


Roll – Up Summary

- Generally, if there are k dimensions, it is possible to define more than 2^k roll - up SQL queries of the type:
 SELECT <attribute_list> AGG(attribute_name)
 FROM <list_of_tables>
 WHERE <list_of_conditions>
 GROUP BY <grouping_list>
- There are exactly 2^k dimensional roll – ups
- The number of hierarchical roll – ups depends on the number of levels of dimension attribute hierarchies
- A simple (with no parallel branches) attribute hierarchy with n levels, gives rise to n roll-ups


Slice_and_Dice Operations

- ***Slice*** operation corresponds to reducing the number of dimensions by taking a projection of facts on a proper subset of dimensions for some selected attribute values of dimensions that are being dropped
- ***Slice*** operation amounts to equality select condition
- ***Dice*** operation amounts to range select condition on one dimension, or to equality select condition on more than one dimension
- Both terms come from visualizing the effect of these operations on a hypercube data set
- Here, the term **select** relates to the **σ** operator of the relational algebra

A Sales Hypercube

	January	February	March	
Eraser	123	99	183	
Linear	23	1	13	
Pencil	1230	1111	2303	

Slice Operation Visualization



Dice Operation Visualization

	January	February
Eraser	123	99
Linear	23	1
Pencil	1230	1111

$$\pi_{ProdId, CityId, MonthId, Amount} (\sigma_{MonthId \in \{January, February\}}(Sales))$$

SQL and Slice_and_Dice Operations

- Consider Sales fact table with *Product*, *City*, and *Time* as dimensions:

Sales (ProdId, CityId, MonthId, Amount)

- A slice operation can be described in the following way

```
SELECT ProdId, CityId, Amount
FROM Sales
WHERE MonthId = 'January' ;
```

- A dice operation would be:

```
SELECT ProdId, CityId, MonthId, Amount
FROM Sales
WHERE MonthId = 'January' OR MonthId = 'February' ;
```

or

```
SELECT CityId, Amount
FROM Sales
WHERE MonthId = 'January' AND ProdId = 'Pencil' ;
```

Pivoting

- In a fact table, each tuple corresponds to at least one measure value and there is a column for each of dimensions
- The simplest view of **pivoting** is that it selects two dimensions to aggregate one of measures
- The aggregated values are often displayed in a grid where each point in the (x, y) coordinate system corresponds to an aggregated value of the measure
- The x and y coordinate values are the values of the selected two dimensions
- The result of pivoting is also called **cross – tabulation**

Extension of a Star Schema

Location

<i>CityId</i>	<i>City</i>
1	Well
2	Nels
3	Auck

SalesPerson

<i>PerId</i>	<i>Name</i>
1	John
2	Susan
3	James
4	Susan
5	Ann

Sales

<i>CityId</i>	<i>PerId</i>	<i>TimId</i>	<i>Amnt</i>
1	1	1	230
1	1	2	300
1	1	8	310
1	2	7	50
2	3	1	550
2	3	5	100
3	4	6	880
3	5	1	60
3	5	2	60
3	5	4	140

Time

<i>TimId</i>	<i>Day</i>
1	Mon
2	Tue
3	Wed
4	Thu
5	Fri
6	Sat
7	Sun
8	Mon

Pivoting on City and Day

	Mon	Tue	Wed	Thu	Fri	Sat	San	SubTotal
Auckland	60	60	0	140	0	880	0	1140
Nelson	550	0	0	0	100	0	0	650
Wellington	540	300	0	0	0	0	50	890
SubTotal	1150	360	0	140	100	880	50	2680

Expressing Pivoting by SQL (body)

```
SELECT City, Day, SUM(Amnt) AS Sales
FROM Sales NATURAL JOIN Location NATURAL JOIN Time
GROUP BY City, Day;
```

<i>City</i>	<i>Day</i>	<i>Sales</i>
Auckland	Monday	60
Auckland	Tuesday	60
Auckland	Thursday	140
Auckland	Saturday	880
Nelson	Monday	550
Nelson	Friday	100

<i>City</i>	<i>Day</i>	<i>Sales</i>
Wellington	Monday	540
Wellington	Tuesday	300
Wellington	Sunday	50

**But, this is only the body
of the pivoting table**

Expressing Pivoting by SQL (subs&total)

```
SELECT Day, SUM(Amnt) AS Total
FROM Sales NATURAL JOIN Time
GROUP BY Day;
```

Day	Total
Mon	1150
Tue	360
Thu	140
Fri	100
Sat	880
San	50

```
SELECT City, SUM(Amnt) AS Total
FROM Sales NATURAL JOIN Location
GROUP BY City;
```


City	Total
Auck	1140
Nels	650
Well	890

```
SELECT SUM(Amnt)
FROM Sales;
```

SUM(Amnt)
2680

Summary of the Cross Tab Example

SELECT City, Day, SUM(Amnt)... GROUP BY City, Day



SELECT SUM(Amnt) FROM Sales

Structures like that one are often called ***cube lattice***

Replacing Pivot Operation by SQL

- To express a pivot operation involving two dimensions, we would have to issue the following SQL/92 statement:

Replacing Pivot Operation by SQL

```
(SELECT City, Day, SUM(Amnt) FROM Sales NATURAL JOIN  
Location NATURAL JOIN Time  
GROUP BY City, Day)
```

```
UNION ALL
```

```
(SELECT City, ' ', SUM(Amnt) FROM Sales NATURAL JOIN Location  
GROUP BY City)
```

```
UNION ALL
```

```
(SELECT ' ', Day, SUM(Amnt) FROM Sales NATURAL JOIN Time  
GROUP BY Day)
```

```
UNION ALL
```

```
(SELECT ' ', ' ', SUM(Amnt) FROM Sales);
```

A Relational Representation of Pivoting

City	Day	Sales
Auckland	Monday	60
Auckland	Tuesday	60
Auckland	Thursday	140
Auckland	Saturday	880
Nelson	Monday	550
Nelson	Friday	100
Wellington	Monday	540
Wellington	Tuesday	300
Wellington	Sunday	50

City	Day	Sales
Auckland		1140
Nelson		650
Wellington		890
	Monday	1150
	Tuesday	360
	Thursday	140
	Friday	100
	Saturday	880
	Sunday	50
		2680

Summary of Basic OLAP Operations

OLAP operation	Use of attribute hierarchy	No of dimensions	Use of aggregate functions	Use of σ	May be replaced by
Roll-up	Y	Same or smaller	Y	N	SELECT... AGG GROUP BY
Drill-down	Y	Same or greater	May use	N	SELECT... AGG GROUP BY
Slice	N	Smaller	N	Y	SELECT... WHERE cond
Dice	N	May be Smaller	N	Y	SELECT... WHERE cond
Pivot / cross tab	Y	Smaller	Y (on many Subsets)	N	Multiple roll-up