

VICTORIA UNIVERSITY OF WELLINGTON  
*Te Whare Wananga o te Upoko o te Ika a Maui*



# ***Materialized Views***

*Lecturer : Dr. Pavle Mogin*

SWEN 432  
*Advanced Database Design and  
Implementation*

# ***Plan for Materialized Views***

---

- Views, OLAP and Data Warehouse
- A motivating example
- Query rewriting as a motivating force for using materialized views
- View materialization
- Examples
  - *Reading :*
    - *Chaudhuri, Dayal : An Overview of Datawarehousing and OLAP Technologies*

# Virtual Views

---

- A SQL view is a virtual table whose content is computed on demand, every time from beginning
- The syntax is:  

```
CREATE VIEW <view_name> [<view_attr_list>]  
AS SELECT ...
```
- Once defined, view can be used in queries, or in defining new views as any other table
- Note: a SQL view is a **virtual** view, it is not materialized

# ***Materialized Views and OLAP***

---

- Views are frequently used in decision support systems to allow data analyst to consider just his/hers part of the business
- Decision support queries are typically aggregate queries over very large fact tables
- Fact tables have a constant content over a relatively long period of time
- To allow fast answers, view materialization is a viable alternative

# ***Views and the Data Warehouse***

---

- A Data Warehouse itself is a (materialized) view of the operational databases and external data sources
- Fast execution of decision support queries is frequently accomplished by materializing views on DW base tables

# Using Virtual Views in Queries

---

- Consider the star schema  
*{Sales (ProdId, TimeId, LocId, Sales\_Data),  
Product (ProdId, Prod\_Name, Category, Price),  
Location (ShopId, Shop\_Name, City, District),  
Time (TimeId, Date, Week, Month, Quarter, Year)}*
- And a virtual view:

```
CREATE VIEW Prod_Shop_Sales  
(ProdId, Category, ShopId, City, TimeId, Sales  
_Data) AS SELECT ProdId, Category, ShopId, City,  
TimeId, Sales_Data  
FROM Product NATURAL JOIN Location NATURAL JOIN  
Sales;
```

# Using Virtual Views in Queries

---

- Consider a query on *Prod\_Shop\_Sales* view (that is still not materialized) :

```
SELECT Category, City, SUM(Sales_Data)
FROM Prod_Shop_Sales
GROUP BY Category, City;
```

- It is very likely that the query processor will evaluate this query by replacing view *Prod\_Shop\_Sales* by its definition

# Query Modification Using A Virtual View

---

```
SELECT Category, City, SUM(Sales_Data)
FROM (
  SELECT p.ProdId, Category, l.ShopId,
  City, TimeId, Sales_Data
  FROM Product NATURAL JOIN Location
  NATURAL JOIN Sales s) AS view
GROUP BY Category, City;
```



# View Materialization and Queries

---

- To speed up the repetitive execution of the previous query, the view *Prod\_Shop\_Sales* can be evaluated and stored
- This way doing two long lasting joins will be saved, every time you execute a query requiring join of:
  - Sales,
  - Product, and
  - Location

tables, projected on:

*{ProdId, Category, ShopId, City, TimeId, Sales\_Data }*

or some of it's proper subsets

# Query Rewriting

---

- Suppose now the view

Prod\_Shop\_Sales

is **materialized**

- A clever query optimizer would rewrite the query

```
SELECT Category, City, SUM(Sales_Data)
FROM Product NATURAL JOIN Location NATURAL JOIN
Sales GROUP BY Category, City;
```

in the following way:

```
SELECT Category, City, SUM(Sales_Data)
FROM Prod_Shop_Sales GROUP BY Category, City;
```

and achieve a considerable saving

# View Materialization and Queries

---

- The query Q:

```
SELECT Category, City, SUM(Sales_Data)
FROM Product NATURAL JOIN Location
NATURAL JOIN Sales NATURAL JOIN Time
WHERE Year = 2000
GROUP BY Category, City;
```

can not be evaluated using **only** the materialized view *Prod\_Shop\_Sales*, since the view does not contain information about year

- But, the materialized view *Prod\_Shop\_Sales* can be joined back to *Time* in the rewritten query and thus be used to answer the query Q

# Creating a Materialized View

- In a conventional Relational Database Management Engine:

```
CREATE TABLE <mat_view_name>...  
INSERT INTO <mat_view_name> ...
```

or

```
CREATE TABLE <mat_view_name>  
AS  
SELECT...
```

- In a Relational On Line Analytical Processing (ROLAP) Engine:

```
CREATE MATERIALIZED VIEW <mat_view_name>  
AS  
(SELECT...) ;
```

# Choosing Views to Materialize

---

- The choice of views to materialize is complex, because the range of views that can be used for a query evaluation is very broad
- On the other hand, materialized views strongly influence the storage occupancy and DW maintenance time
- When deciding which views to materialize, one should consider the following issues:
  - How many queries potentially can be sped up,
  - How much space will be required to store the views, and
  - How will the views influence the DW maintenance (update)
- It is a goal to materialize a small, carefully chosen set of views that can be used to evaluate the majority of important queries

# Two Symmetric Queries

---

- Consider the following two symmetric queries:

```
SELECT Category, SUM(Sales_Data)
FROM Product NATURAL JOIN Sales
GROUP BY Category;
```

```
SELECT City, SUM(Sales_Data)
FROM Location NATURAL JOIN Sales s
GROUP BY City;
```

# Speeding Up Two Symmetric Queries

---

- There exist a number of possible ways to speed up these queries:
  - To precompute and store as materialized views both joins (*Sales with Product*, and *Sales with Location* ),
  - To precompute and store as materialized views both queries in their entirety, or
  - To define and materialize the following view:

```
CREATE MATERIALIZED VIEW Prod_Loc_Sales (Category,  
City, Total) AS
```

```
SELECT Category, City, SUM(Sales_Data)
```

```
FROM Sales NATURAL JOIN Product NATURAL JOIN Location  
GROUP BY Category, City;
```

- There also exist some other ways

## Question for You

---

- Suppose a clever query optimizer
- The view *Prod\_Loc\_Sales* (given on the previous slide) favors one of the two queries

```
SELECT Category, SUM(Total)
FROM Prod_Loc_Sales
GROUP BY Category;
```

```
SELECT City, SUM(Total)
FROM Prod_Loc_Sales
GROUP BY City;
```

- Which one and why?



# ***Materialized View Advisor***

---

- ROLAP Engines offer automatic advise to a DW DBA on which materialized views to build and which ones to drop
- This advising is done using statistical data gathered by Data Warehouse querying in a previous time period

# Summary

---

- OLAP queries are typically aggregate queries over very large fact tables
- To allow fast answers, view materialization is a viable alternative
- Materialized views strongly influence the storage occupancy and DW maintenance time
- It is the goal to materialize a small, carefully chosen set of views that can be used to evaluate a majority of the most important queries