

VICTORIA UNIVERSITY OF WELLINGTON
Te Whare Wananga o te Upoko o te Ika a Maui



MongoDB

Distributed Write and Read

Lecturer : Dr. Pavle Mogin

SWEN 432
*Advanced Database Design and
Implementation*

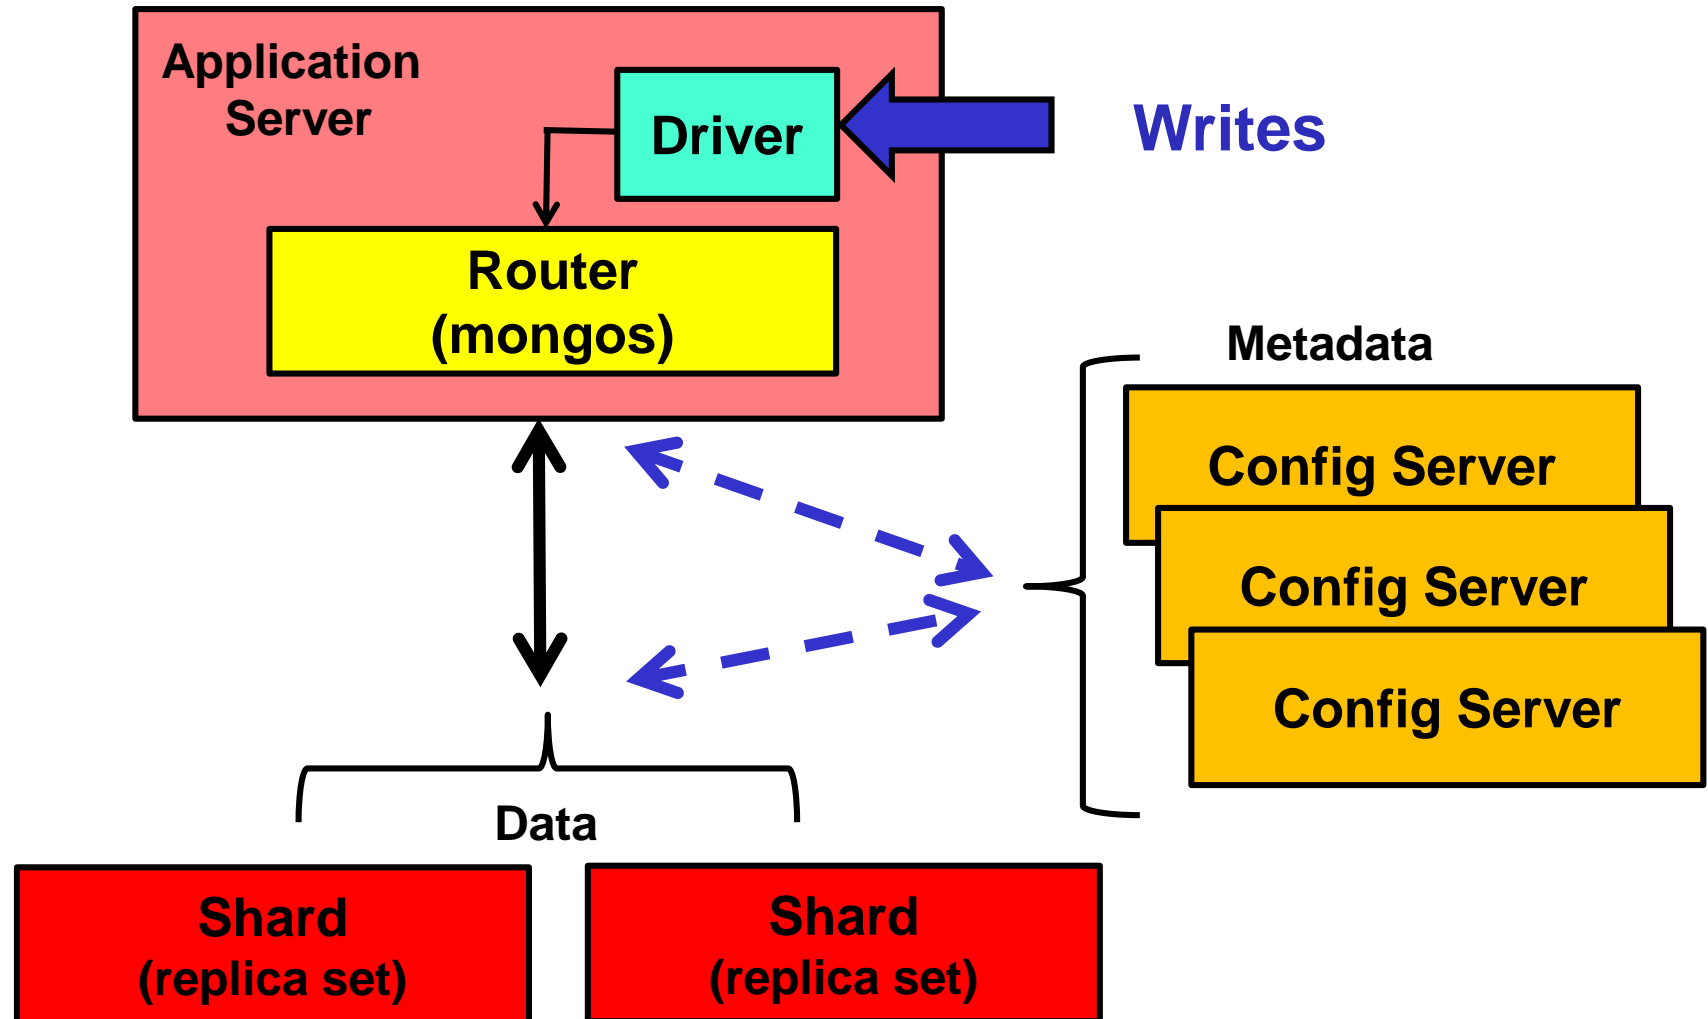
Plan for Distributed Write and Read

- Distributed Write
 - Write on Sharded Cluster
 - Write on Replica Sets
- Write Concern
 - Bulk () Method
- Distributed Queries
- MongoDB and Transaction Processing
 - ***Readings:***
 - *Have a look at Readings on the Home Page*

Write Operations on Sharded Clusters

- For sharded collections in a sharded cluster, the `mongos` directs write operations from applications to shards that are responsible for the portion of the data set using the sharding key value
- The `mongos` gets needed metadata information from the config database residing on config servers

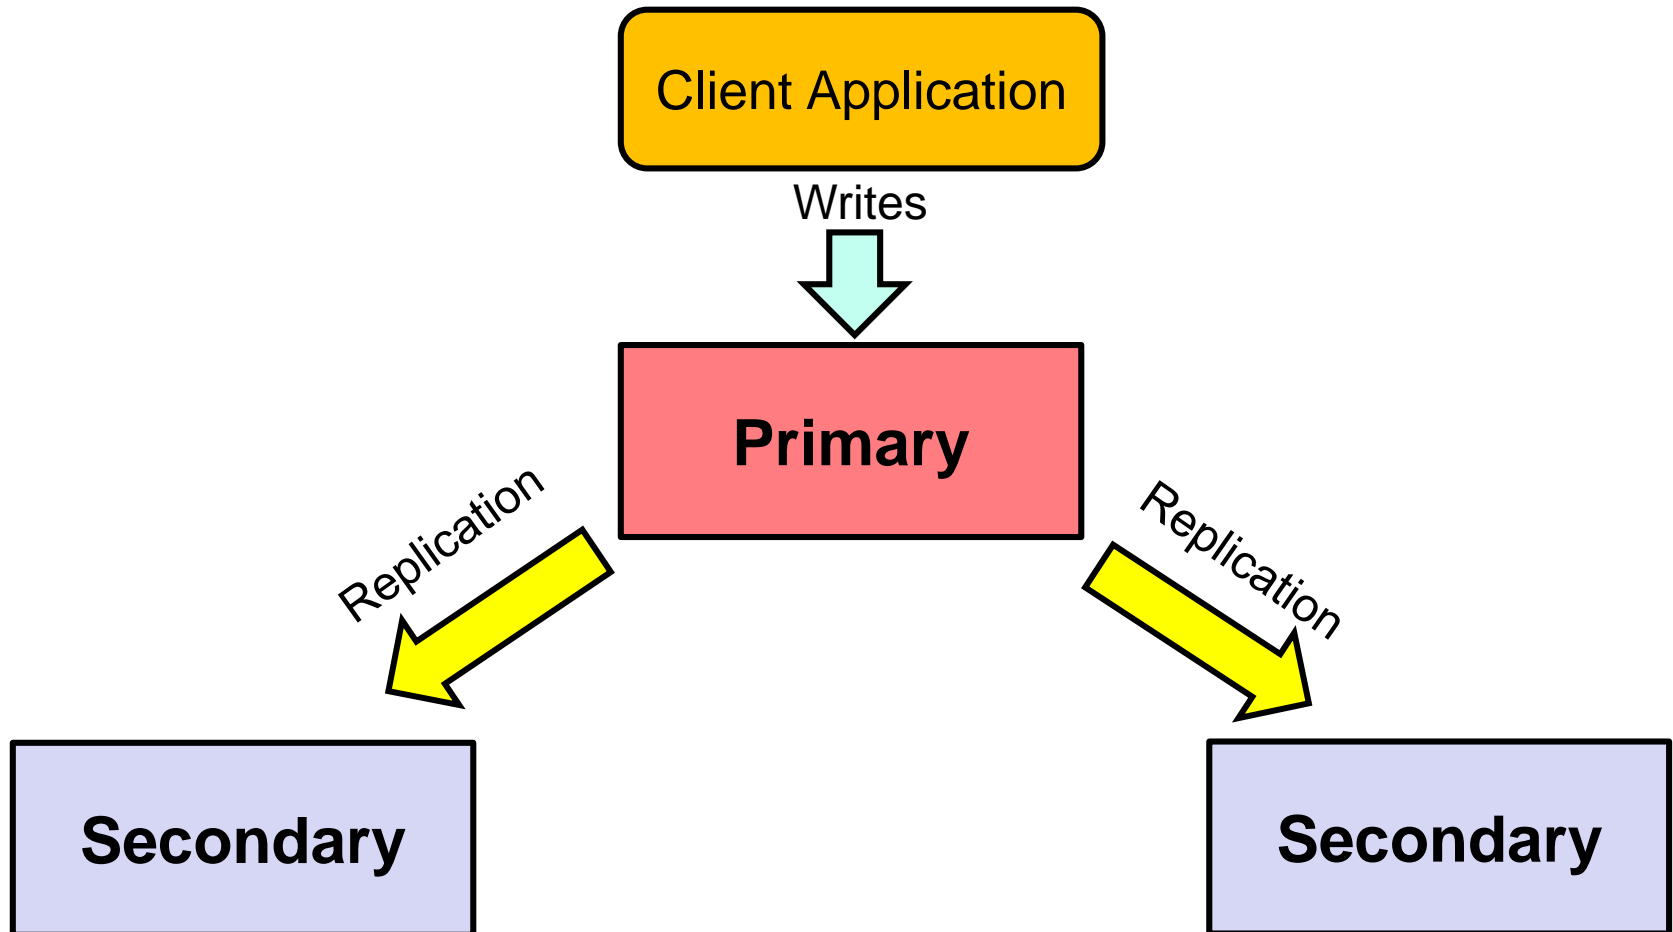
Sharded Cluster



Write Operations on Replica Sets

- In replica sets, all write operations go to the set's primary
- The primary applies the write operations and then records the operations on its operation log (oplog)
 - Oplog is a reproducible sequence of operations to the data set
- Secondary members of the set continuously replicate the oplog by applying operations to themselves in an asynchronous process

Replica Set Operations



Write Concern

(1)

- *Write concern* describes the guarantee that MongoDB provides when reporting on the success of a write operation
- The strength of the write concerns determine the level of guarantee
- When inserts, updates and deletes have a *weak* write concern, write operations return quickly
- In some failure cases, write operations issued with weak write concerns may not persist
- With *stronger* write concerns, clients wait after sending a write operation for MongoDB to confirm the write operations

Write Concern

(2)

- MongoDB (**version 2.6**) provides different levels of write concern:
 - Unacknowledged (lowest level),
 - Acknowledged (default),
 - Journaled, and
 - Replica Acknowledged (highest level)
- Clients may adjust write concern to ensure that the most important operations persist successfully to an entire MongoDB deployment
- For other less critical operations, clients can adjust the write concern to ensure faster performance rather than ensure persistence to the entire deployment

Insert Multiple Documents with Bulk ()

1. Initialize a Bulk () operator for the collection

```
var bulk =  
    db.myclasses.initializeUnorderedBulkOp();
```

2. Add a number of insert operations to the bulk object using bulk.insert () method

```
bulk.insert(doc1);
```

```
...
```

```
bulk.insert(docn);
```

3. Execute the execute () method on the bulk object

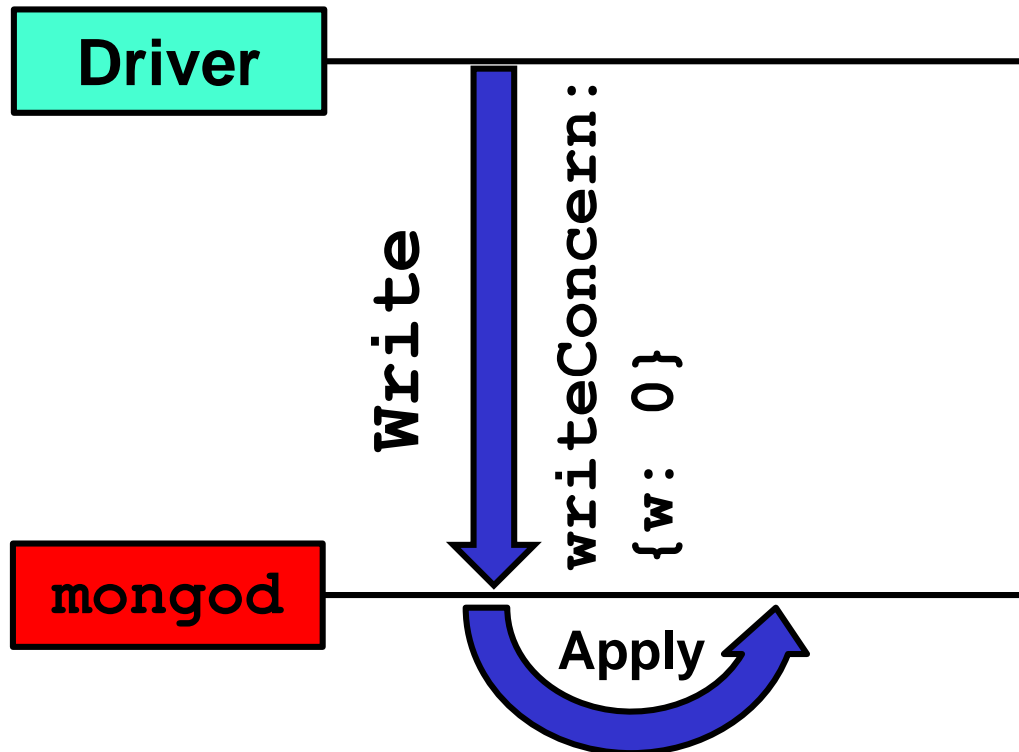
```
bulk.execute({w: "j"});
```

- The execute () method has an optional parameter **w** for specifying the write concern level

4. The method returns a BulkWriteResult

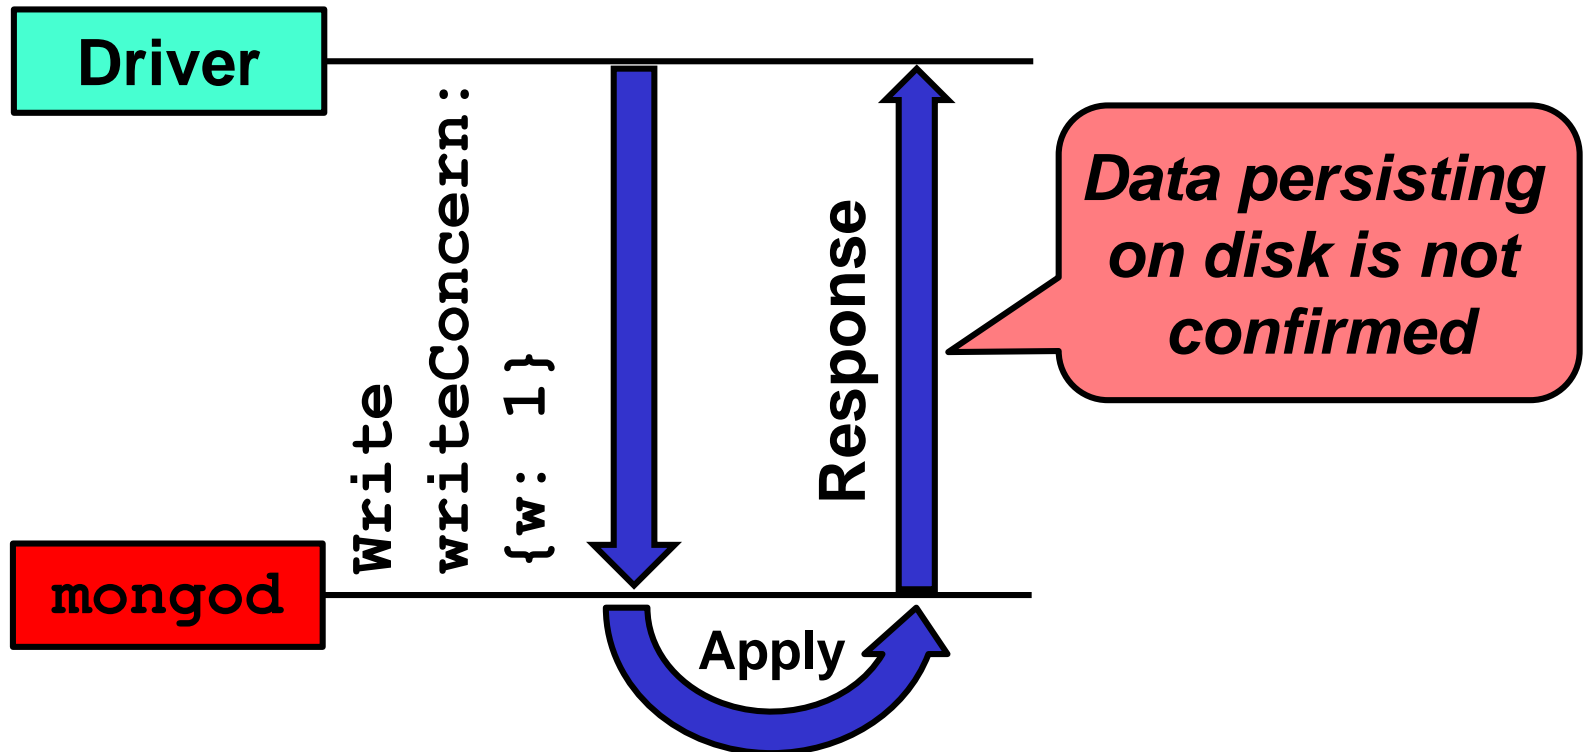
Write Concern: Unacknowledged

- If **{w: 0}**, MongoDB does not acknowledge the receipt of a write operation



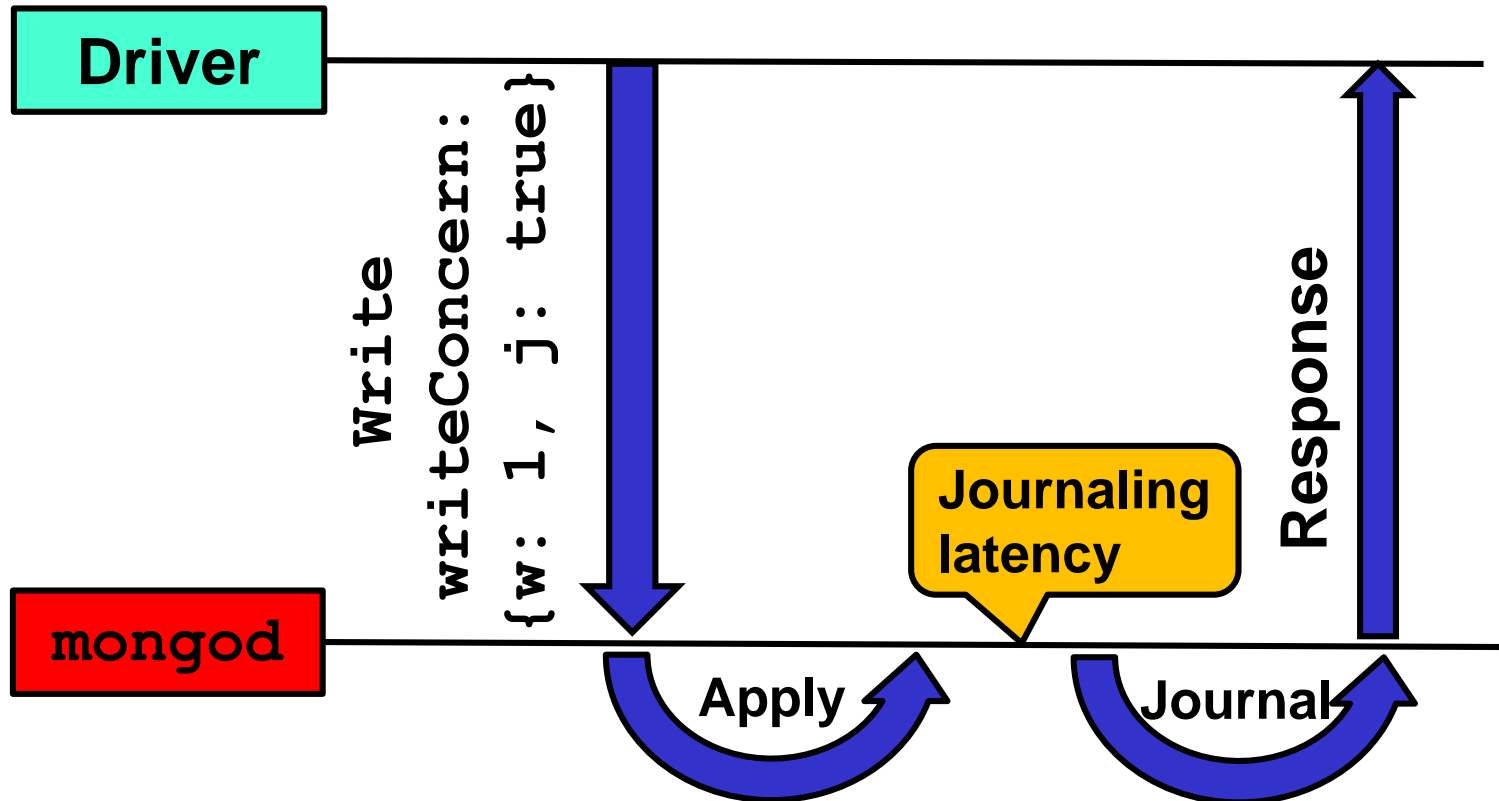
Write Concern: Acknowledged

- If **{w: 1}**, MongoDB confirms that it applied a change to the in-memory data



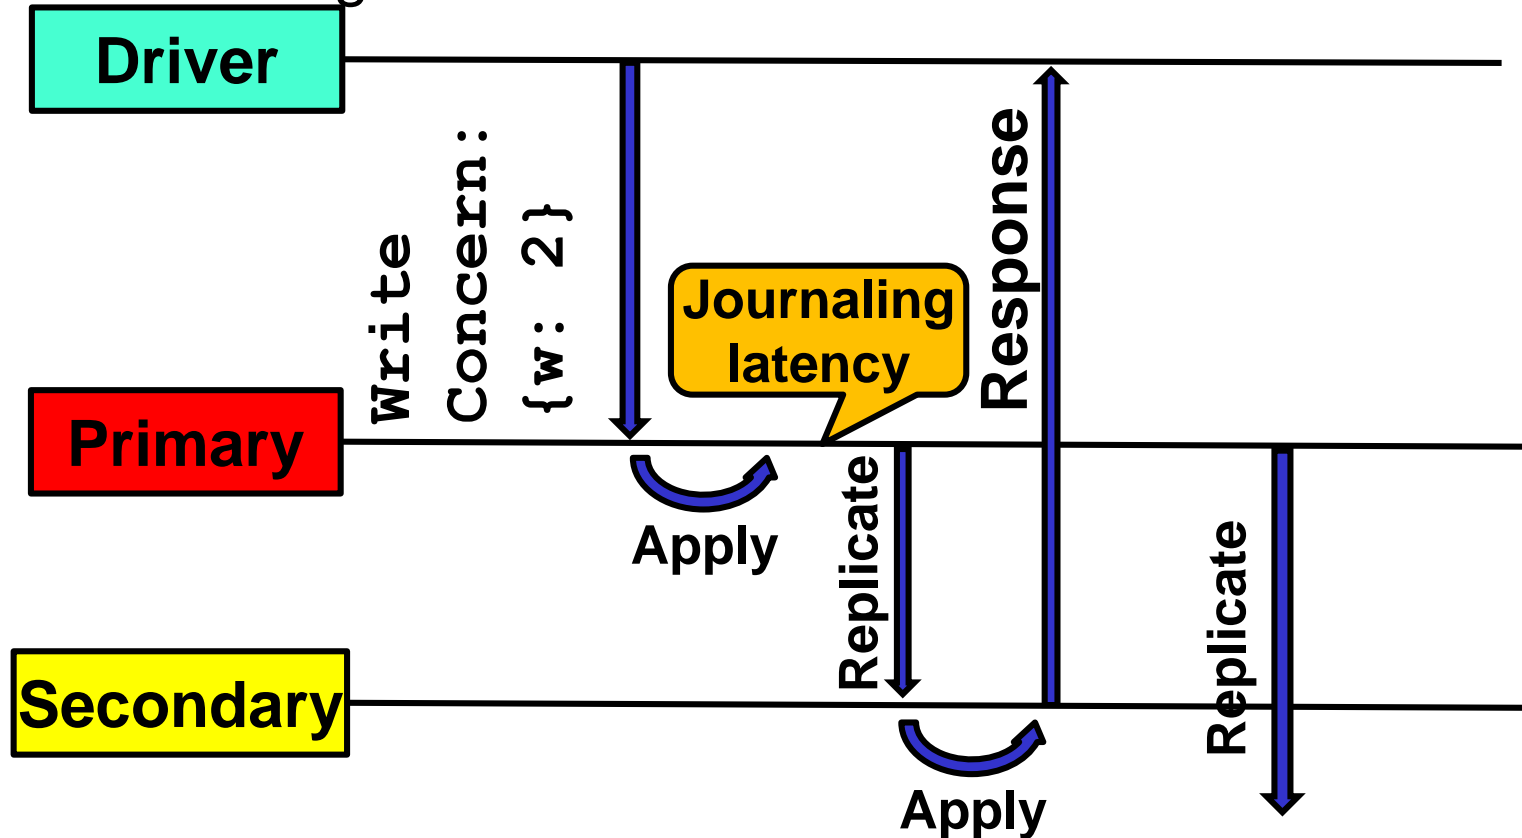
Write Concern: Journalled

- If `{w: 1, j: true}`, MongoDB confirms that it committed data on (master's) disk



Write Concern: Replica Acknowledged

- If **{w: 2}**, the first secondary to finish in memory application of primary's oplog operation, returns acknowledgment



Distributed Queries

- Applications issue operations to one of `mongos` instances of a sharded cluster
- Read operations are most efficient when a query includes the collection's shard key
 - Otherwise the `mongos` must direct the query to all shards in the cluster (scatter gather query) and that might be inefficient
- By default, MongoDB always reads data from a replica set's primary

Reading From a Secondary

- Reading from a secondary server is possible and justified if there is a need :
 - To balance the work load,
 - To allow reads during failover, but
 - Eventual consistency can be guaranteed, only
- To allow reading from a slave server, one of the following set-ups are needed:
 - Modifying the read preference mode in the driver, which results in a permanent change, or
 - Connecting to a slave server shell and issuing the following commands :

```
db.getMongo().setSlaveOk()
```

```
use <db_name>
```

```
db.collection.find()
```

Read Isolation

- MongoDB allows clients to read documents inserted or modified before committing modifications to disk, regardless of write concern level
 - MongoDB performs journaling frequently, but only after a defined time interval
- If the `mongod` terminates before the journal commits, even if a write returns successfully, queries may have read data that will not exist after the `mongod` restarts
 - This is a *read uncommitted* transaction anomaly.
- When `mongod` returns a successful *journalled write concern* (`"j"`), the data is fully committed to disk and will be available after `mongod` restarts

Atomicity

- A write operation is atomic on the level of a single document, even if the operation modifies multiple embedded documents *within* a single document
- When a single write operation modifies multiple documents, the modification of each document is atomic, but the operation as a whole is not atomic and other operations may interleave
 - There exists the `$isolated` operator that can *isolate* a single write operation
 - But it does **not** work on sharded clusters

Transaction Like Semantics

- Since a single document can contain multiple embedded documents, single-document atomicity is sufficient for many practical use cases
- For cases where a sequence of write operations must operate as if in a single transaction, a *two-phase commit* can be implemented in an application
- However, the two-phase commit can only offer transaction-*like* semantics
- Using two-phase commit ensures data consistency, but it is possible for applications to return intermediate data during the two-phase commit or rollback

Concurrency Control

- In relational databases, concurrency control allows multiple applications to run concurrently without causing data inconsistency or conflicts
- MongoDB does not offer such mechanisms
- Instead, there are techniques to avoid some sorts of inconsistencies:
 - Unique indexes used with certain methods like `findAndModify()` prevent duplicate insertions or updates
 - Also, there are certain programming patterns that can be applied to avoid concurrency control anomalies, like the lost update anomaly

Summary

- Routers direct client read and write operations to shards and their replica sets using meta data from config servers
- All writes go to the master server
- By default, all reads also go to the master server
- Write Concern is the guarantee that MongoDB provides when reporting on the success of a write operation
 - Weak write concern: fast, but not very reliable
 - Strong write concern: slower, but more reliable
- By default, queries are of the type “read uncommitted”
- Queries based on the shard key value are the fastest
- Transaction like behavior is achievable to some extent